

Announced the NVIDIA L40S GPU — a universal data center processor designed to accelerate the most compute-intensive applications — available from leading server makers in a broad range of platforms, including NVIDIA OVX™ and NVIDIA AI-ready servers with NVIDIA BlueField® DPUs, beginning this quarter. Unveiled NVIDIA MGX™, a server reference design available this quarter that lets system makers quickly and cost-effectively build more than 100 server variations for AI, HPC and NVIDIA Omniverse™ applications. Announced NVIDIA Spectrum-X™, an accelerated networking platform designed to improve the performance and efficiency of Ethernet-based AI clouds, which is shipping this quarter. Joined with global system makers to announce new NVIDIA RTX™ workstations with up to four new NVIDIA RTX 6000 Ada GPUs, as well as NVIDIA AI Enterprise and NVIDIA Omniverse Enterprise software, expected to ship this quarter. Launched general availability of cloud instances based on NVIDIA H100 Tensor Core GPUs with Amazon Web Services, Microsoft Azure and regional cloud service providers. Partnered with a range of companies on AI initiatives, including: ServiceNow and Accenture to develop AI Lighthouse, a first-of-its-kind program to fast-track the development and adoption of enterprise generative AI capabilities. VMware to extend the companies' strategic partnership to ready enterprises running VMware's cloud infrastructure for the era of generative AI with VMware Private AI Foundation with NVIDIA. Snowflake to provide businesses with an accelerated path to create customized generative AI applications using their own proprietary data. WPP to develop a generative AI-enabled content engine that lets creative teams produce high-quality commercial content faster, more efficiently and at scale while staying fully aligned with a client's brand. SoftBank to create a platform for generative AI and 5G/6G applications based on the GH200, which SoftBank plans to roll out at new, distributed AI data centers across Japan. Hugging Face to give developers access to NVIDIA DGX™ Cloud AI supercomputing within the Hugging Face platform to train and tune advanced AI models. Announced NVIDIA AI Workbench, an easy-to-use toolkit allowing developers to quickly create, test and customize pretrained generative AI models on a PC or workstation and then scale them, as well as NVIDIA AI Enterprise 4.0, the latest version of its enterprise software. Set records in the latest MLPerf training benchmarks with H100 GPUs, excelling in a new measure for generative AI.